

Unit 10

Interconnection busses

Course 2020/2021



Contents

■ Busses

- The bus concept
- Technological aspects
- Topologies
- Interconnecting busses
- Bus hierarchy

■ Current busses

- Trends
- PCI and PCIe
- SATA
- USB and Firewire
- Today's bus hierarchy

■ Transfers within the computer

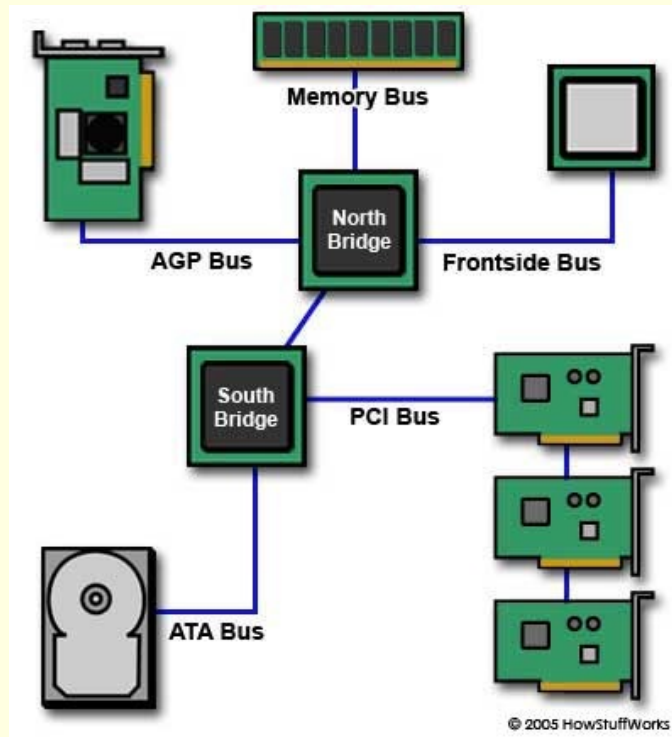
- The roles of the bus controller and OS
- Examples of timing



Busses

The need for interconnection

- CPU, memories and I/O devices have varying bandwidth requirements
- They are interconnected with different types of busses



The bus concept

- A bus is a means for communicating two or more devices, enabling:
 - Addressing: selecting one of the connected devices and the addressable elements within them
 - Synchronization: a means of signalling that a device is ready to operate
 - Transfer: the effective transmission of data among devices
- Other, optional functions are:
 - Power supply
 - Arbitration
 - Hot plug capability ...
- **Bus cycle:** time interval taken to complete an elementary data transfer between interconnected devices:

Technological aspects

■ Bus requirements and specifications

- **Bandwidth:** bus speed – must be enough for supporting the working speed of the connected devices
- **Length**
 - Some devices are close to each other, within few cm. (CPU, memory controller, graphics adapter, etc.)...
 - ...whereas others may be farther apart (printer, scanner, network devices, etc.) A flexible cable is needed to connect them to CPU and memory
- **Standardization**
 - Interchangeable devices (disks, graphics adapter, keyboard, etc.) need to accommodate the specification of standard busses
 - Busses interconnecting *faster* devices (CPU, system clock, memory controller, etc.) are more hardware-dependent and they need not be subject to a standard specification

Technological aspects

■ Electrical issues

- **Electrical noise:** Computer components and neighbouring equipment produce electromagnetic interference
 - The problem grows with cable length, and is reduced with shielding
- **Degradation and clock skew:** signals lose shape and synchronicity when they travel the bus cables. The problem grows with:
 - Faster bus transfers (shorter bus cycles)
 - Cable bending, that alters the geometry and electrical characteristics of cables in the bus
- **Crosstalk:** lines in a bus interfere each other. The problem grows:
 - With the number of wires in the bus
 - With insufficient shielding (e.g. when it is sacrificed for cable flexibility)

Technological aspects


■ Physical aspects

- A bus is composed by a set of electric wires
 - Often times, a *shielding* conductor wraps the others to reduce electromagnetic interference
 - Wires are spaced from each other to reduce crosstalk
 - Cable length is limited
- A bus specification includes detailed description of the allowed connectors
- Each bus uses a **protocol**: a particular way of commanding signals in the bus, including also error detection
- A bus may be **serial** or **parallel**, according to the number of data bits transferred in a single *bus cycle*

Technological aspects

■ Parallel busses

- All word bits are transferred simultaneously in a single bus clock cycle
- Parallel bus example (unidirectional)




Implementation using two registers with parallel load

1. Transmitter writes its register
2. Data arrives in the receiver register and is written there
3. Receiver can then read the transmitted data from its own register

Technological aspects

■ Serial bus

- Bits are transmitted serially. For n -bit words, n elementary transmissions (bus clock cycles) are needed to transfer a full word
- A serial bus requires only one wire for data
 - Few more lines for signalling, power, etc.



Implementation with two shift registers

1. Transmitter writes the output register
2. Word is transferred bit by bit from the transmitter's serial output to the receiver's serial input
3. Receiver can read the whole word only at the end of transfer

Technological aspects

■ Serial and parallel busses: comparative

- The control unit for the serial connection is more complex
- Parallel cables are heavier and more rigid. Connectors are more fragile and less convenient
- Under ideal conditions, parallel is faster, but...
 - ..."ideal" means no noise and perfect wires (no capacitance, no inductance)
- For high frequencies (\sim GHz), parallel busses can only be very short (\sim cm) due to clock skew and cross-talk

Type	Control complexity	Electrical issues	No. of wires and pins
Parallel	lower	serious	many
Serial	higher	simple	few

Technological aspects


■ Bus bandwidth

- Let f denote the bus clock frequency
- For a parallel bus of width w data bits, the bandwidth is:
$$B = f \times w/8 \text{ Bps}$$
 - Examples:
 - PCI-X. $f = 133.3 \text{ MHz}$, $w = 64 \text{ bits}$, $B = 1066.6 \text{ MBps}$
 - Parallel ATA-133. $f = 66 \text{ MHz}$, $w = 16 \text{ bits}$, $B = 133 \text{ MBps}$
- For a serial bus, $B = f \text{ bps}$, BUT redundant bits for synchronization and error control must be subtracted to obtain the *effective* bandwidth (e.g., 8b/10b, 128b/130b, 128b/132b,...)
 - Examples
 - PCIe-1x (version 2). $f = 5 \text{ GHz}$, using 10 bits/byte $\rightarrow B = 500 \text{ MBps}$
 - SATA-3 Gbps. $f = 3 \text{ GHz}$, using 10 bits/byte $\rightarrow B = 300 \text{ MBps}$

Bus topologies


■ **Multipoint** (or *multidrop*)

- Limited number of addressable devices
- Examples: ATA (limit = 2), PCI (limit = 2^{32} or 2^{64})



■ **Point to point**


- Direct connection between just two devices
 - No need for device selection
- Examples: RS-232, AGP




Bus topologies

■ Star busses

- In their simplest form, point to point




- Extended with additional switching elements (*hubs*)
- Example: USB, SATA




Bus topologies

■ Daisy chain

- Point-to-point connection, with transducers (two connections)
 - If address on the bus is not *mine*, pass it on to next in chain
- Examples: SCSI, Firewire




- With hubs: tree topology (Firewire)



Interconnecting busses


- Interconnecting two distinct busses requires solving two issues:
 - **Physical matching:** each bus has its own specs for signals (frequencies, voltages...) Signals must be adapted from one bus to the other
 - **Logical matching:** Devices in one bus must be addressable from devices in the other



Interconnecting busses

■ Bridges


- Bridges help preserving logical unity
 - They give the view of a single addressing space. Programs do not care about which bus a peripheral is physically connected to
- A bridge is seen as another device on each bus. It provides access to devices on the other bus
 - On bus A, the addresses for P3 and P4 select the bridge
 - On bus B, the addresses for P1 and P2 select the bridge
- When selected, the bridge translates the signals to the other bus



Interconnecting busses

■ Bus adapters


- Adapters offer an interface to access devices on a different bus
- The interconnected busses have separate addressing spaces
- Programs must select the adapter and program its registers in order to access devices in the other bus through the adapter
- Example: P1, P2 and adapter share the same addressing space on bus A. P3 and P4 can be selected from bus A by writing their addresses in the adapter



Interconnecting busses

■ Bandwidth


- When two busses are interconnected by a bridge, the maximum bandwidth is:
 - The bandwidth of the used bus, when the transfer does not cross the bridge
 - The minimum bandwidth in the path, when the transfer needs to cross the bridge
- Examples: $T1 \rightarrow B_A$ $T2 \rightarrow \min(B_A, B_B)$ $T3 \rightarrow B_B$



Bus hierarchy

■ The *system controller*


- Combination of DRAM controller plus a set of bridges connecting:
 - **The processor bus.** Proprietary design, depends on the particular CPU. Targets maximum bandwidth
 - **The main memory bus.** Complies with a given standard (e.g., DDR3)
 - **The standard expansion bus.** Targets maximum compatibility. Provides standard connectors for compatible peripheral adapters



Bus hierarchy

■ The system bus


- A system bus connects the devices that are mapped in the addressing space of the processor
- It is typically formed by the processor bus, the system controller and the expansion bus
 - The expansion bus may connect to other busses via bridges



Bus hierarchy

■ I/O busses (or peripheral busses)

- A set of standard I/O busses connected to the system bus via bus adapters
- Each I/O bus has its own addressing space. Programs must use the bus adapter interface in order to access the peripherals



Bus hierarchy

■ Summary

- Typical computers include a set of busses
 - Busses are interconnected by bridges and adapters
 - Each bus is chosen to satisfy certain criteria: compatibility, bandwidth, etc.
- The closer to CPU-memory, the faster the bus



Current busses

Trends

■ Years 1980...2000 (approx.):

- Parallel, multipoint expansion busses (ISA - PCI, NuBus...)
- Parallel peripheral busses (Centronics, SCSI, ATA) for higher bandwidth demands (scanner, hard disk, etc.)
- Serial peripheral busses (RS-232) for slower devices (mouse, keyboard) or remote devices (modems, printers)

■ Years 2000...2005

- Parallel expansion busses (PCI and AGP)
- Parallel peripheral busses only for internal discs (ATA)
- Serial peripheral busses (USB and Firewire)

■ Currently (2005...)

- Serial expansion busses (PCI express)
- Serial peripheral busses (SATA, USB and Firewire)
- Only the processor bus is still mostly parallel (not clear what the future will be in this regard...)

Trends

Devices' bandwidth

Device	MBps (approx.)
processor (Core Duo 2GHz)	10000
SDRAM channel DDR3 400 MHz	6400
Graphics display (1600 × 1200, 50 fps)	300
Hard disk (7200 rpm, 1600 sectors/track)	100
DVD (20x)	27 (1x = 1.385 MBps)
CD-ROM (52x)	7.8 (1x = 153.6 KBps)

Peripheral busses

Peripheral	old	current
mouse, keyboard	RS232	USB
Graphics display	PCI, AGP	PCIe
Internal hard disk	ATA	SATA
External hard disk	SCSI	USB, Firewire, eSATA
Optical drive	ATA	SATA


PCI

■ Features

- Multipoint parallel bus designed to operate as system bus
- Introduced in 1993, still in use for compatibility

■ Bandwidth

- In its many versions, bandwidth has evolved from the initial 133 MBps to 4 GBps
 - PCI 2.3 (conventional):
533 MBps (64bits/66MHz)
 - PCI-X 1.0: 1GBps (64bits/133MHz)
 - PCI-X 2.0: 2GBps (64bits/266MHz)
4GBps (64bits/533MHz)



PCI-express (PCIe)


- Evolved version of the classical PCI as system bus
- Serial bus, using 8b/10b (128b/130b in PCIe 3.0, 4.0, 5.0)
- Point-to-point bus, characterized by the number of *lanes* L_B (1x, 2x, 4x, 8x, 12x, 16x or 32x). Each lane enables serial transfers at:
 - 250 MB/s – 2.5 GT/s (version 1.0, 1.1) [8b/10b]
 - 500 MB/s – 5 GT/s (v. 2.0)
 - 985 MB/s – 8 GT/s (v. 3.0) [128b/130b]
 - 1969 MB/s – 16 GT/s (v. 4.0)



PCI-express (PCIe)


■ PCIe peripheral adapters

- Each adapter has a given nr. L_p of connections (1x, 2x, etc...) to bus lanes
- At start-up, the system checks the number of available lanes to talk with the peripheral through the bus: $\min\{L_B, L_P\}$
- Common values for L_p : 8x for the graphics adapter, 1x for audio
- Examples with PCIe 2.0:



PCI-express

■ PCI-express adapters and connectors



16x adapter



1x adapter



Peripheral busses



■ SATA bus


- Most frequently used with storage units: hard disks and optical drives
- Serial connection (maximum length of 1 m), 8b/10b encoding
- Two topologies:
 - Point-to-point: one peripheral per bus
 - Star, with one level of *multipliers* (~hubs). Up to 15 peripherals
- Versions and bandwidth:
 - SATA 1.5 Gb/s: 150 MBps (effective)
 - SATA 3 Gb/s "SATA II": 300 MBps
 - SATA 6 Gb/s: 600 MBps
- A version (*eSATA*) exists for external peripherals, for cables up to 2 m and bandwidth 3 Gbps (300 MBps)



Peripheral busses




- SATA Express bus
 - Supports SATA and PCIe 3.0
 - Up to 1969 MBps (2 PCIe lanes)
 - SATA-compatible connector
 - Conceived for solid state disks (SSDs)

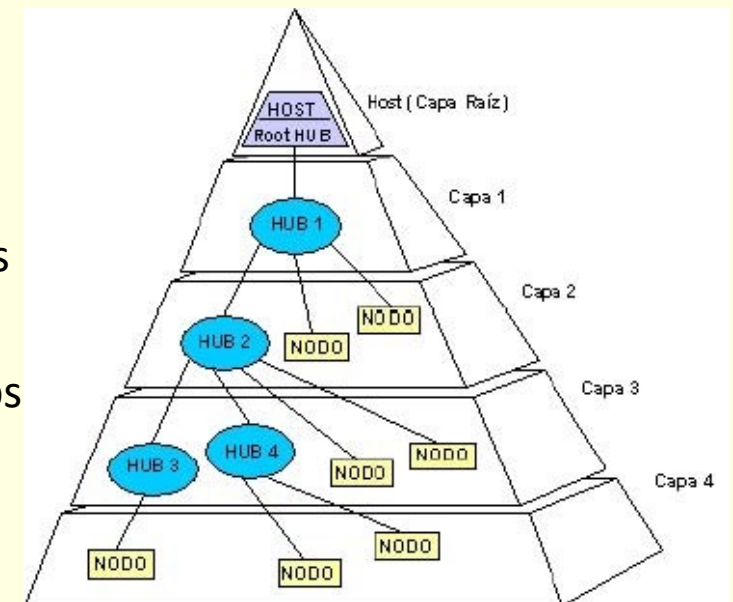


Peripheral busses

■ USB bus



- General purpose peripheral bus with serial transfer, star topology
- One bus controller, several hubs where peripherals are connected
- Cables:
 - Maximum length = 5 m (w/o power)
 - Power lines included (5 V, 0.5 A typ.)
- Up to 6 levels of hubs
- Up to 127 devices per bus
- Versions and bandwidth:
 - 1.0 - 12 Mbps (8b/10b) → 1.5 MBps
 - 2.0 - 480 Mbps (8b/10b) → 60 MBps
 - 3.0 - 4.8 Gbps (8b/10b) → 480 MBps
 - 3.1 - 10 Gbps (128b/132b) → 1.2 GBps
 - 3.2 - 20 Gbps (128b/132b) → 2.4 GBps



Peripheral busses



■ Firewire (IEEE 1394, i.Link)

- General purpose peripheral bus
 - Up to 63 peripherals
- Serial transfer, daisy chain topology
 - Maximum length: 4,5 m for one cable, 72 m the whole bus
- Very versatile
 - Allows connection between computers
 - Allows direct communication between devices connected to the bus
 - Widely used for professional video
- Versions and bandwidth:
 - Firewire 400 Mbps
 - Firewire 800 Mbps
 - Firewire 1600 and 3200 Mbps

Peripheral busses

- M.2 connectors
 - ✓ Enable multiple interfaces via a 75-pin connector

KEY	INTERFACES	COMMON USES
A	PCIe x2, USB 2.0, I2C, DisplayPort x4	Wi-Fi/Bluetooth, cellular cards
B	PCIe x2, SATA, USB 2.0, USB 3.0, audio, PCM, IUM, SSIC, I2C	SATA and PCIe x2 SSDs
E	PCIe x2, USB 2.0, I2C, SDIO, UART, PCM	Wi-Fi/Bluetooth, cellular cards
M	PCIe x4, SATA	PCIe x4 SSDs

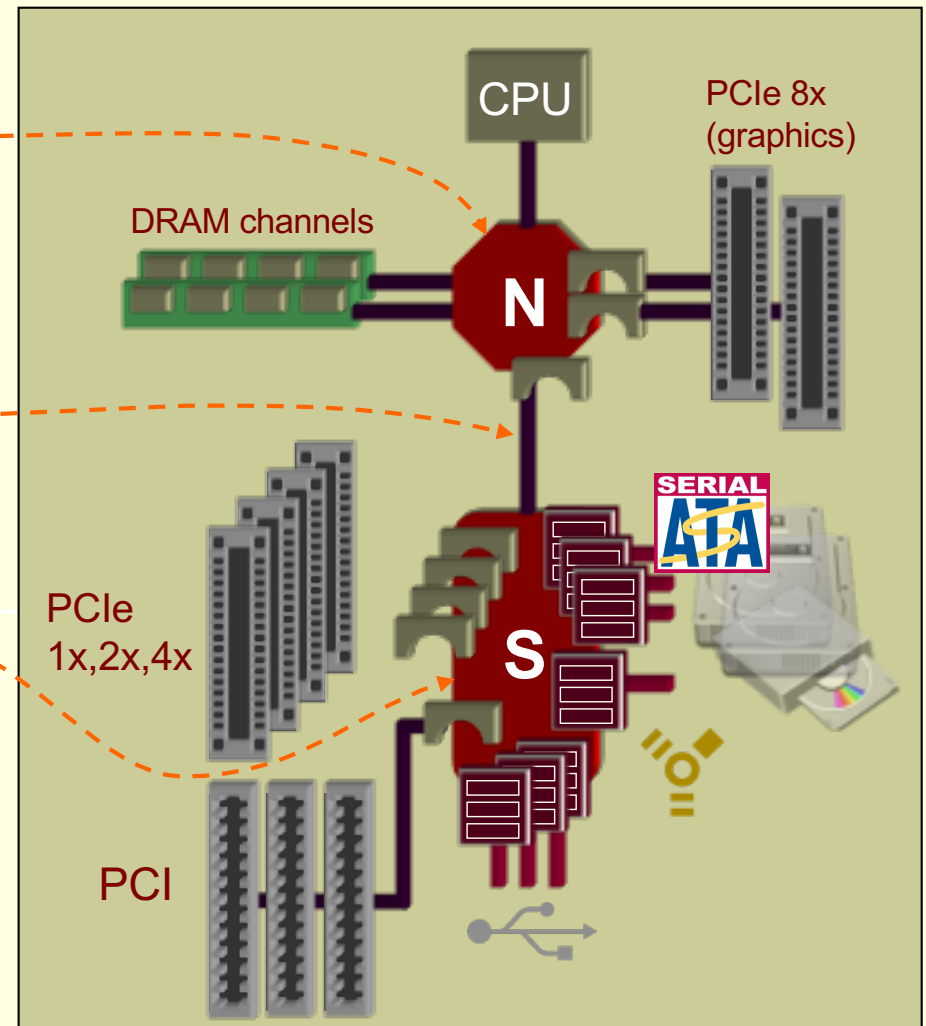


Two B- and M-keyed SSDs (left)
M-keyed SSD (right)

Trends


■ Current bus hierarchy

- The system controller (*northbridge*) controls access to the fastest busses: 2 DRAM channels and video adapter
- A (usually proprietary) bus connects the *northbridge* with the *southbridge*
- The *southbridge*, *system hub* or I/O controller is a collection of bridges and I/O bus adapters





Motherboard connections

ASUS P5Q PRO



Motherboard connections




Trends

■ External units

■ Via bus adapter:

- Combine a disk unit and a bus adapter (e.g. *external USB disk drive*)
- The included bus adapter makes the translation between the general-purpose I/O bus (USB, Firewire) and a specific bus (e.g. SATA)
- The applicable IODTR* is limited by the slowest bus (the I/O bus; e.g. USB is slower than SATA)



■ Without bus adapter

- Sometimes a faster connection is directly available for external units (e.g. eSATA)
- The applicable IODTR is high, greater than the SDTR **

*IODTR = Input-Output Data Transfer Rate, a parameter of the bus

**SDTR = Sustained Data Transfer Rate, a parameter of the device




Transfers within the computer

Transfers within the computer

■ Data flow within the computer

- Main memory is the *central* resource
- All traffic crosses the system controller
 - The CPU reads/writes data and fetches instructions from there (when cache misses)
 - DMA block devices make transfers Main Memory ↔ Peripheral
 - PIO devices make transfers Main Memory ↔ CPU ↔ Peripheral



Transfers within the computer

■ Data flow control

- The system bus and many I/O busses support concurrent transfers
 - Elements connected to the same bus compete for it
 - Controllers and bus arbiters handle the multiplexing of concurrent transfers
 - On each bus, the bandwidth consumption is the sum of bandwidths consumed by all transfers taking place
 - Utilization of the bus cannot surpass the maximum bus bandwidth. The bus controller/arbiter can limit the speed of individual transfers to keep the transit within the bus limits
- When a transfer traverses several busses, the effective bandwidth is limited by the slowest bus involved

Transfers within the computer

■ The role of the OS

- Regular programs use OS services to use the peripherals
 - Especially the file-system-related functions
- The OS services program the peripheral adapters and carry the corresponding PIO or DMA transfers
 - With file-system services, the OS performs additional housekeeping functions (directory modifications, file table updating, etc.)
- Under ideal circumstances, the **theoretical bandwidth** depends only on the speeds of the busses (IODTR) and peripherals (SDTR) involved
- In actual fact, the **effective bandwidth** is reduced by other factors (time taken to program the devices, arbitration conflicts, etc.)

Transfers within the computer


■ Real-time aspects

- Transfers with no real-time restrictions (*as fast as possible*)
 - File transfers, file read/write, internet browsing...
 - Transfers occur at the maximum available bandwidth, taking a time $T = (\text{amount of data})/B$
- Transfers with real-time restrictions (*within given deadlines*)
 - Typically, multimedia: audio/video playing/recording, streaming...
 - They must comply with some given real-time restriction (frames per second, audio samples per second)
 - If the available bandwidth is sufficient, transfers occur at the appropriate speed. Otherwise, the results will be defective
 - A special case are critical real-time applications, where failing to meet the timing requirements can lead to catastrophic results
 - e.g. control systems in aviation, air traffic control, medical equipment, satellites, power stations, etc.

Examples

■ Example 1

- A program reads a full 1 GB file from a hard disk connected to a SATA bus (1.5 Gbps)
 - Disk: SDTR = 100 MBps; DMA transfer
 - Context:




What we can conclude

- The maximum bandwidth is limited by the hard disk
- Minimum time for transfer is 10 s (1 GB / 100 MBps)
- Bandwidth consumption in the busses:
 $100/12800 = 0.78\%$ (M); $100/2000 = 5\%$ (NS); $100/150 = 67\%$ (SATA)

Examples

■ Example 2

- Full read of a 1 GB file in a external disk, connected via USB 2.0
 - Disk: SDTR = 100 MBps; DMA transfer
 - Context:




What we can conclude:

- Bandwidth limited by the USB bus
- Minimum time for transfer is 16.7 s (1 GB / 60 MBps)
- Bandwidth consumption in the busses:
 $60/12800 = 0.47\%$ (M); $60/2000 = 3\%$ (NS); 100% (USB)

Examples

■ Example 3

- Copy a 1 GB file from disk (connected to SATA 1.5 Gbps) to another disk (connected to bus USB 2.0)
- Disks: 100 MBps (SDTR); DMA transfer
- Context:




What we can conclude:

- Bandwidth limited by USB bus
- Minimum transfer time is 16.7 s (1 GB / 60 MBps)
- Bandwidth consumption in the busses:
 $2 \times 60 / 12800 = 0.94\%$ (M); $2 \times 60 / 2000 = 6\%$ (NS); $60 / 150 = 40\%$ (SATA); 100% (USB)

Examples

■ Example 4

- Copy a 1 GB file to the same disk (connected to SATA 1.5 Gbps)
- Disk: 100 MBps (SDTR); DMA transfer
- Context:




What we can conclude:


- Bandwidth limited by the disk
- Minimum transfer time is 20 s ($2 \times 1\text{GB} / 100 \text{ MBps}$)
- Bandwidth consumption in the busses:
 $100/12800 = 0.78\%$ (M); $100/2000 = 5\%$ (NS); 67% (SATA)

Audio transfers

- Three factors impact size and required bandwidth:
 - Sampling frequency (f_{clk}) – 44.1 / 48 kHz, x2: 88.2 / 96 kHz ...
 - Sample size (**bit depth**) – 8 bits, 12 bits, 16 bits, 24 bits ...
 - Number of **channels** – 1 (mono), 2 (stereo), n+1 (n.1 systems)



Channel systems



Required bandwidth for transmission of uncompressed audio streams:

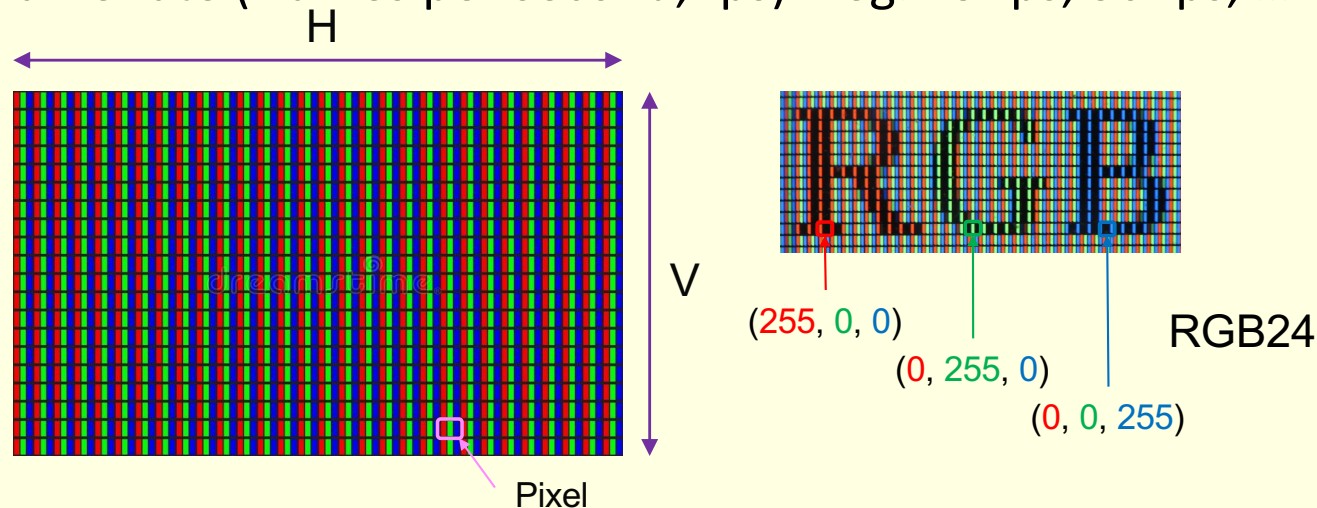
$$B_{audio} = f_{clk} \times \text{bit depth} \times \text{channels} \quad [\text{kbps}]$$

Compressed audio (eg. MP3, compr. rate $\simeq 10:1$), "**coded at f kbps**"

$$B_{comp_audio} = f \quad [\text{kbps}]$$

Video transfers

- Again, three factors impact size and required bandwidth:
 - Size of image ($H \times V$ pixels) – eg. 1280×720 , 1920×1080 , ...
 - Nr. of bits per pixel (bit depth) – eg. 24 bits in RGB24 (1 B/colour)
 - Frame rate (frames per second, fps) – eg. 25 fps, 30 fps, ...



Required bandwidth for transmission of uncompressed video:

$$B_{video} = H \times V \text{ pixels} \times \text{bit depth} \times \text{frame rate} \quad [\text{Mbps}]$$

Compressed video+audio (eg. MPEG-4), "*coded at f Mbps*"

$$B_{comp_audio} = f \quad [\text{Mbps}]$$

Examples


■ Example 5 (with real-time restrictions)

- Playing a DVD movie (video + audio)
 - Transfer contents (mpeg-2) from DVD to Main Memory
 - Meanwhile, the CPU decodes the mpeg-2 to obtain the video frames and PCM audio
 - Video frames need be transferred to the display at the proper frequency
 - Audio needs also be transferred to the sound card
- Context:
 - Movie at 30 fps (frames per second)
 - DVD: mpeg-2 encoding at 10 Mbps
 - Assume the CPU is fast enough to decode mpeg-2 at proper speed
 - Audio 5.1, 16 bits at 48 KHz
 - Frame size is 1600 x 1200 pixels, 24-bit colour

Examples

■ Example 5

■ Context:



The transfer is NOT limited by the busses:

- DVD read: 10 Mbps = 1.25 MBps
- Graphics adapter: $1600 \times 1200 \times 3 \times 30 = 172.8$ MBps
- Audio: $6 \times 48000 \times 2 = 576$ KBps = 0.576 MBps

Busses utilization:

- $1.25/150 = 0.83\%$ (SATA); $172.8/2000 = 8.6\%$ (PCIe-graphics);
 $0.576/250 = 0.23\%$ (PCIe-audio); $(1.25 + 0.576)/2000 = 0.091\%$ (NS);
 $(1.25 + 0.576 + 172.8)/12800 = 0.14\%$ (Memory)

Examples

■ Example 6 (with real-time restrictions)

- Play an uncompressed, high definition silent movie from hard disk (30 fps, 1920 x 1080 pixels, 24-bit colour depth)
 - Each second, read 30 frames of 1920x1080x3 bytes from disk
 - No need to decode frames
 - Each second, write 30 frames in graphics memory
 - The bandwidth required by each transfer is $\approx 187 \text{ MBps}$
- Context:
 - Display: 1920 x 1080 pixels, 24-bit colour
 - Disk: 100 MBps connected to SATA 3Gbps



**The disk is not fast enough
Playback will be defective**