



Unit 2.- Task synchronization



Concurrency and Distributed Systems



Teaching Unit Objectives

- ▶ Review the concept of concurrent programming
 - ▶ Identify the problems related to concurrent programming
 - ▶ Understand the cooperation mechanisms needed to implement a concurrent application.
 - ▶ Mechanisms for communication and synchronization
 - ▶ Identify the parts of the code that can cause problems (critical sections) and how to protect them
- ▶ Know a programming language that supports concurrent programming: Concurrency in Java
 - ▶ Introduce the Java basic mechanisms for supporting concurrent programming



Content

► Task synchronization

▶ Communication mechanisms

- ▶ Shared memory

- ▶ Message exchange

▶ Types of synchronization

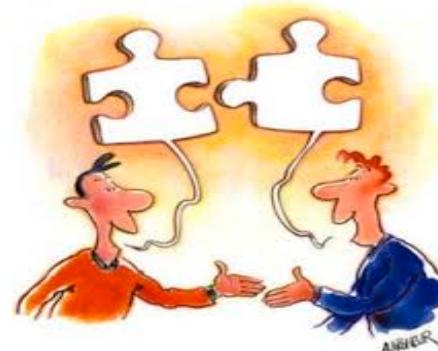
- ▶ Mutual Exclusion

- Critical Section

- Locks

- ▶ Conditional Synchronization

- ▶ Concurrency implies parallelism and cooperation among threads
- ▶ Cooperation requires:
 - ▶ a) communication (information interchange)
 - ▶ b) synchronization (setting certain order in specific cases)





Communication Mechanisms

- ▶ **Communication Mechanisms:**
 - a) **Shared Memory:** threads share space in memory (variables or shared objects)
 - b) **Message Exchange (or Message Passing):** sender/receiver are potentially in disjoint spaces of memory
- ▶ We focus here on **shared memory (option 'a')**
 - ▶ It requires synchronization mechanisms to coordinate tasks
 - ▶ Employed in languages such as Java and C#, and libraries such as *pthreads*
 - ▶ We will see the classical concurrency model
- ▶ We will see *option b* in **Distributed Systems**
 - ▶ The current trend is to implement distributed applications



Shared memory: problems that can appear

- ▶ **Race Conditions:** inconsistent modification of shared memory
 - ▶ Errors may occur when multiple threads access shared data.
 - ▶ Inconsistent views of the shared memory may result in errors.

Synchronization prevents
these problems



Problems with the use of shared memory →

Example 1

- ▶ Example: **c++** can be decomposed into three steps:
 1. Retrieve the current value of **c**
 2. Increment the retrieved value by 1
 3. Store the incremented value back in **c**

(the same for **c--** but decrementing)

- ▶ Important! Even simple statements can be translated to multiple steps by the Java Virtual Machine

```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
  
    public long getCount() {  
        return count;  
    }  
...  
Counter c= new Counter();  
// initialized to 0
```



Example 1

- ▶ EXAMPLE: we have an object `c` of type `Counter` and two threads A and B

- ▶ `c` initially has value 0
 - ▶ A executes `c.add(3)`
 - ▶ B executes `c.add(2)` concurrently

- ▶ The final value of `c` should be 5

- ▶ But the final value depends of the **exact interleaving** (scheduling) when running A and B
 - ▶ There are **thread interferences** when two operations, running in different threads, but acting on the same data, **interleave**.

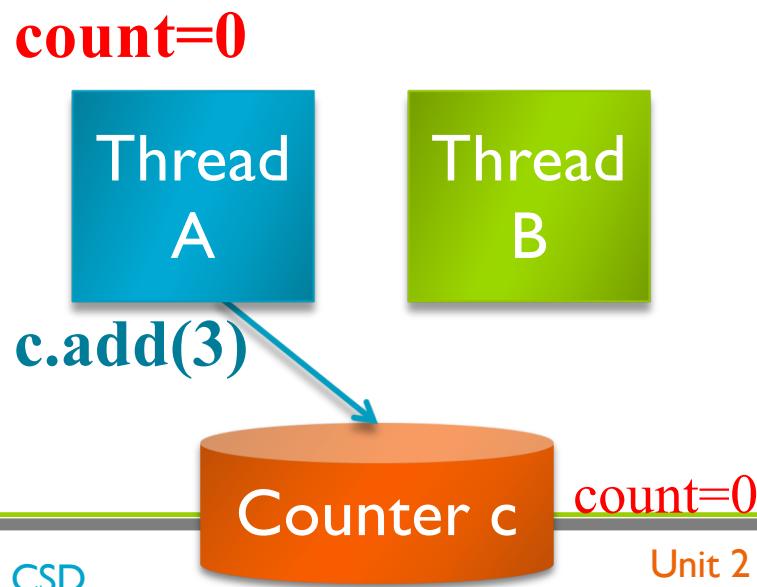
```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
  
    public long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter();  
// initialized to 0
```

Example 1

► Interleaving example:

Thread A executes `c.add(3)`

A loads `c.count` (local copy=0)



```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
    public long getCount() {  
        return count;  
    }  
...  
Counter c= new Counter(); //initialized to 0
```

Example 1

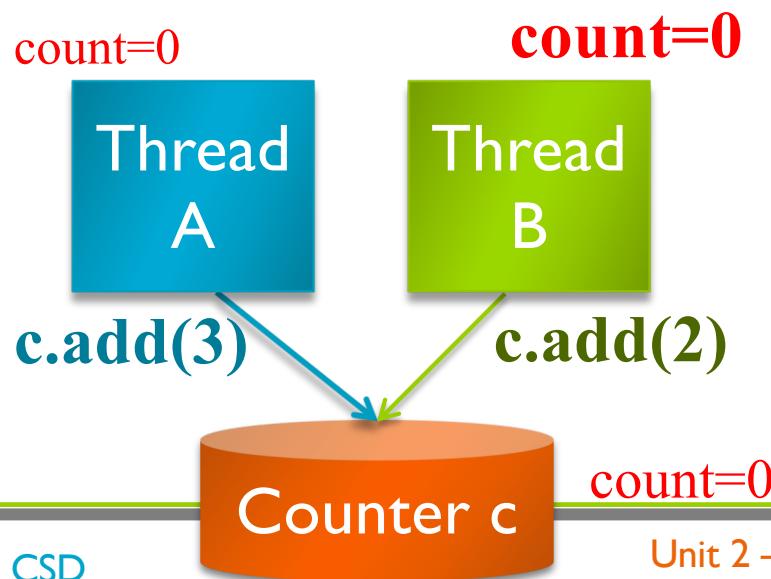
► Interleaving example:

Thread A executes `c.add(3)`

A loads `c.count` (local copy=0)

Thread B executes `c.add(2)`

B loads `c.count` (local copy=0)



```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
    public long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Example 1

► Interleaving example:

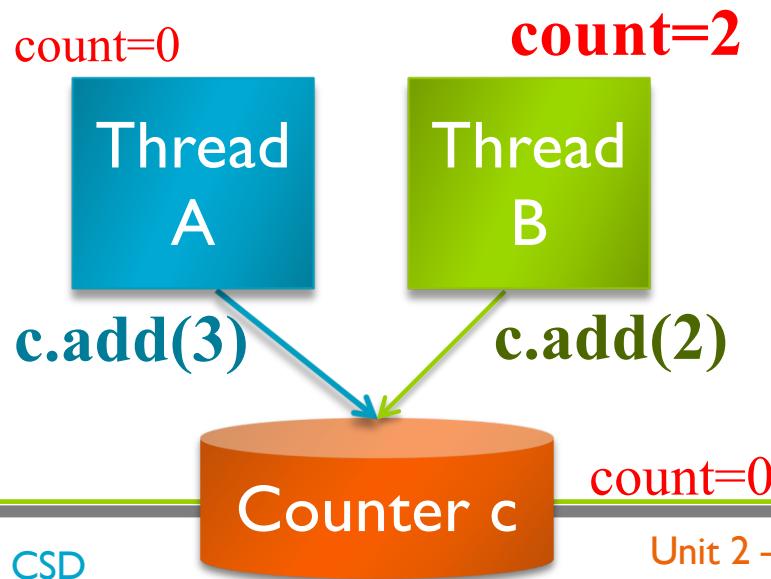
Thread A executes `c.add(3)`

A loads `c.count` (local copy=0)

Thread B executes `c.add(2)`

B loads `c.count` (local copy=0)

B adds 2 (local copy=2)



```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
    public long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Example 1

► Interleaving example:

Thread A executes `c.add(3)`

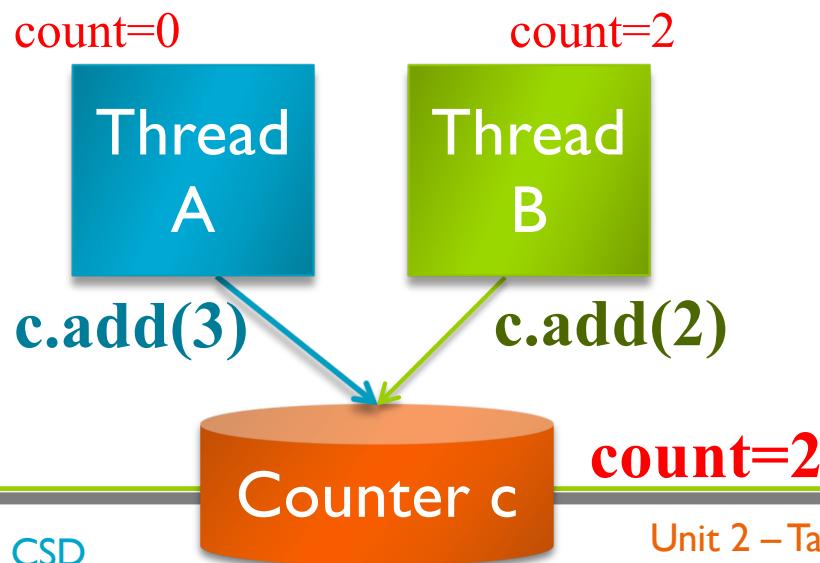
A loads `c.count` (local copy=0)

Thread B executes `c.add(2)`

B loads `c.count` (local copy=0)

B adds 2 (local copy=2)

B updates the heap (`c.count=2`)



```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
  
    public long getCount() {  
        return count;  
    }  
...  
Counter c= new Counter();
```

Example 1

► Interleaving example:

Thread A executes c.add(3)

A loads c.count (local copy=0)

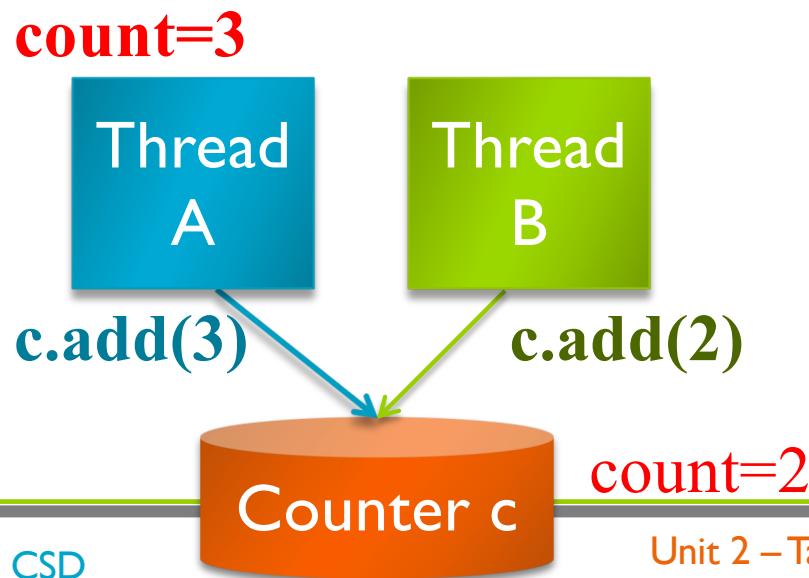
A adds 3 (local copy=3)

Thread B executes c.add(2)

B loads c.count (local copy=0)

B adds 2 (local copy=2)

B updates the heap (c.count=2)



```
public class Counter {  
    protected long count = 0;  
    public void add(long x) {  
        count += x;  
    }  
  
    public long getCount() {  
        return count;  
    }  
...  
Counter c= new Counter();
```

Example 1

► Interleaving example:

Thread A executes c.add(3)

A loads c.count (local copy=0)

A adds 3 (local copy=3)

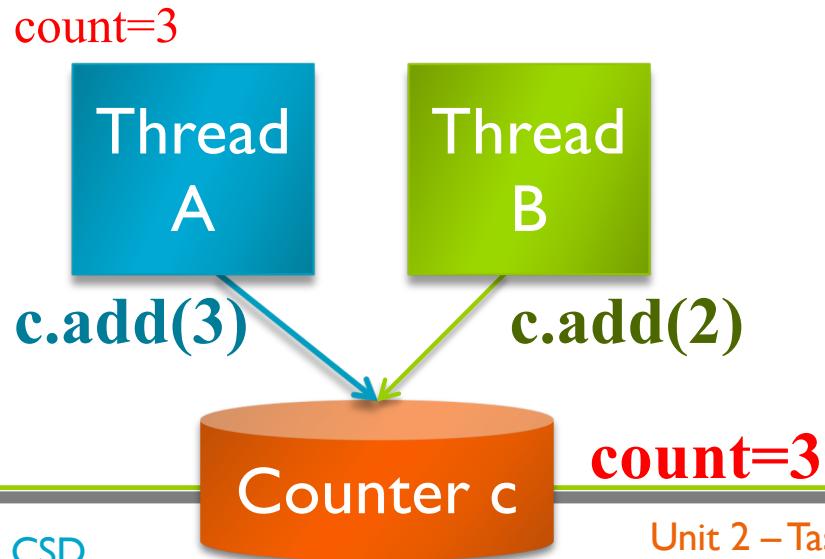
A updates the heap (c.count=3)

Thread B executes c.add(2)

B loads c.count (local copy=0)

B adds 2 (local copy=2)

B updates the heap (c.count=2)



```
public class Counter {
    protected long count = 0;
    public void add(long x) {
        count += x;
    }
    public long getCount() {
        return count;
    }
    ...
}
```

Counter c= new Counter();

Example 1

► Interleaving example:

Thread A executes `c.add(3)`

A loads `c.count` (local copy=0)

A adds 3 (local copy=3)

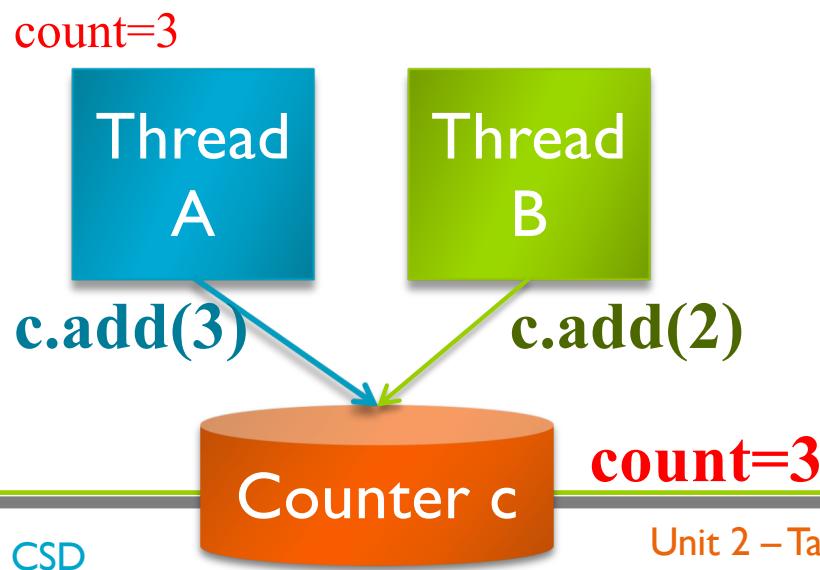
A updates the heap (`c.count=3`)

Thread B executes `c.add(2)`

B loads `c.count` (local copy=0)

B adds 2 (local copy=2)

B updates the heap (`c.count=2`)



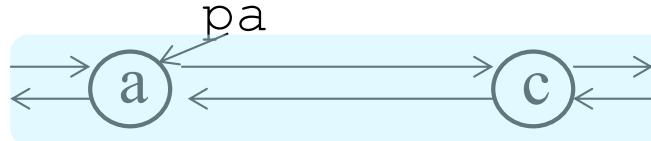
Final value: 3
We expected 5 !!!

There has been a
race condition

Example 2 - Inserts **b** between **a** and **c** in a double-linked list

As atomic action

`Node pa=... //insert after pa`

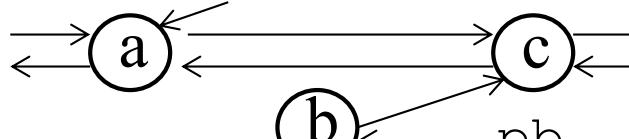


Steady state

`Node pb = new Node(b);`



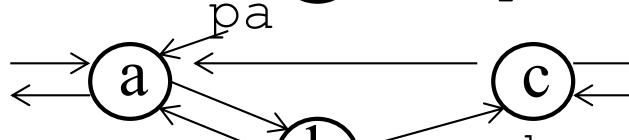
`pb.next = pa.next;`



`pb.prev = pa;`



`pa.next = pb;`



`(pb.next).prev = pb;`



Intermediate States



Content

▶ Task synchronization

▶ Communication mechanisms

- ▶ Shared memory
- ▶ Message exchange

▶ Types of synchronization

▶ Mutual Exclusion

- Critical Section
- Locks

▶ Conditional Synchronization



Goals of synchronization

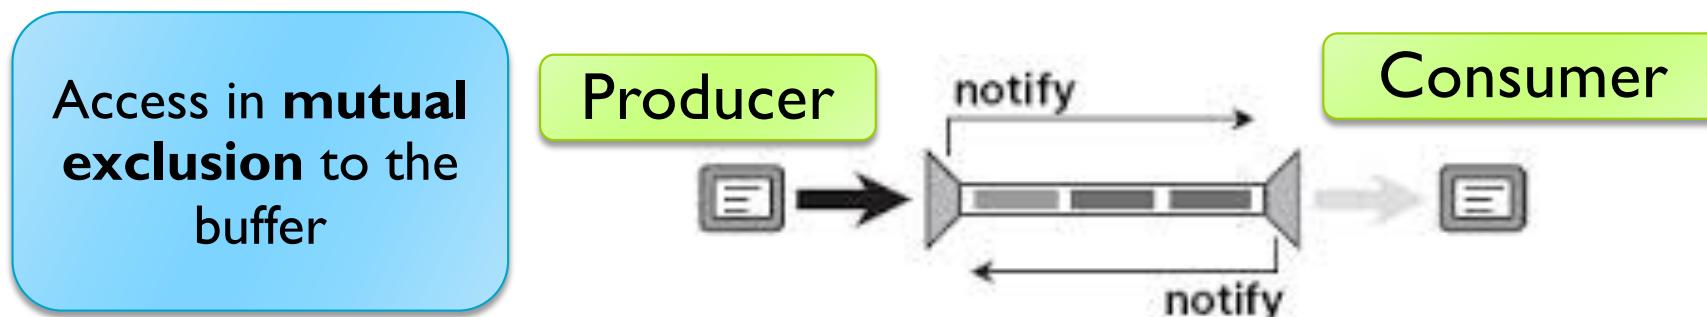
Goals of Synchronization Mechanisms

- Ensure execution order of sentences
- Respect restrictions on code execution

- ▶ **Types of synchronization:**
 - ▶ **Mutual Exclusion**
 - ▶ **Conditional Synchronization**

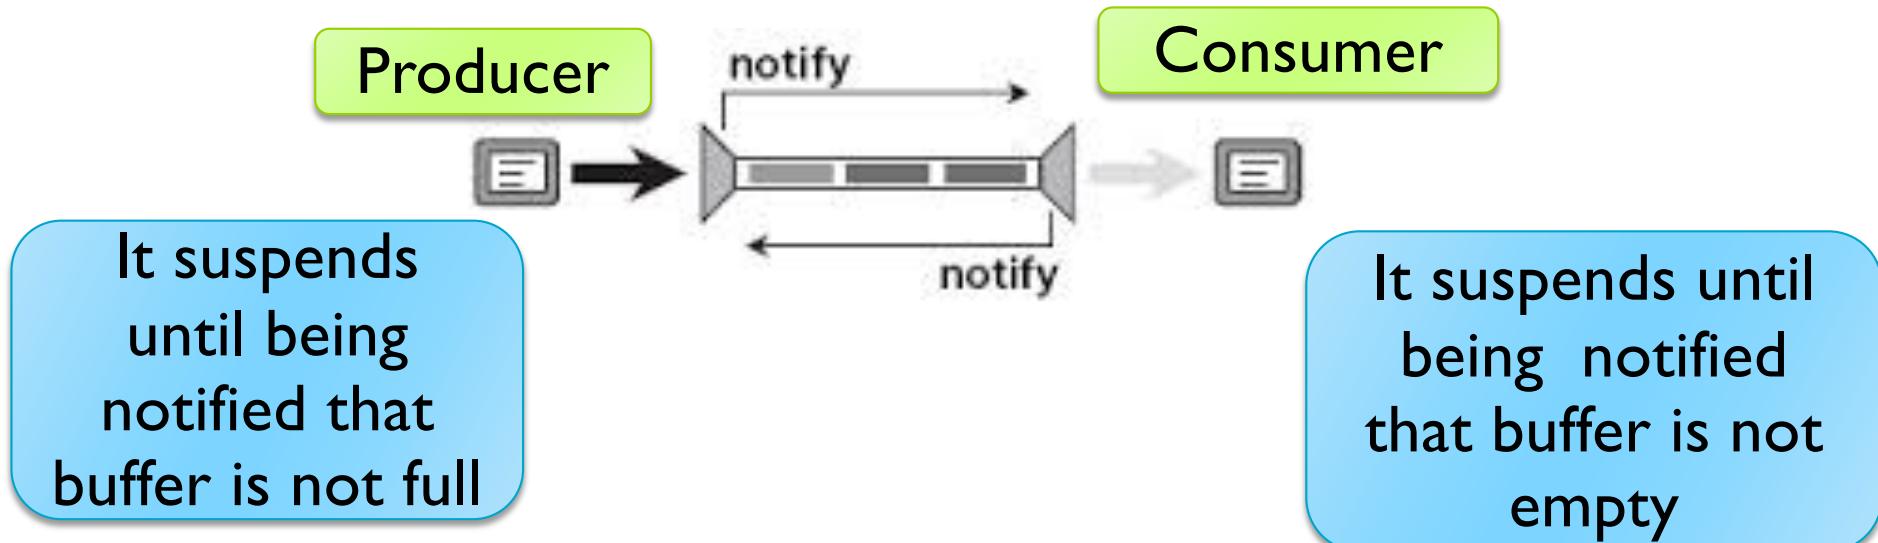
▶ Mutual exclusion

- ▶ A section of code is executed in ***mutual exclusion*** when **only one thread** can run this section at any time.
- ▶ Needed to avoid interferences among threads.
- ▶ Required for shared variables/objects



▶ Conditional Synchronization

- ▶ A thread must delay its execution until a specific condition holds.
- ▶ The condition often depends on the value of some shared variables.
- ▶ Other threads, when modifying these variables, will get the condition to be fulfilled, thus reactivating the suspended threads





Critical Section → What is it?

Critical Section (CS) : fragment of code that can cause race conditions

- ▶ Access to *local variables* → *it is not critical*
- ▶ Access to *immutable shared objects* → not critical
- ▶ Access to ***shared variables/objects*** → **critical**



Critical Section (CS): fragment that accesses to variables or objects shared by more than one thread



Critical Section → How can you protect it?

...
Rest of code

....

input protocol

Critical Section

output protocol

...

Rest of code
(Non-critical Section)

....

Input/Output protocols that ensure (all together):

- **Mutual Exclusion**
- **Progress:** every service requested is completed at some point
- **Bounded Waiting:** there exists a bound, or limit, on the number of times that other processes are allowed to enter their critical sections after a process has **made a request to enter its critical section** and before that request is granted.

Thread-safe code: code that can be run concurrently by different threads safely (i.e. without race-condition problems)



Critical Section → How can you protect it?

- ▶ Solution strategies depending on the level:

Hardware

- Disabling interruptions
- Only for encoding the OS kernel

Operating System

- System calls for using semaphores, mutexes, events, conditions

Programming Language

- *Locks, conditions, monitors, critical sections, protected objects*
- Examples: Concurrent Pascal, Modula-2, Modula-3, Ada, Java, C#

Locks → What is it? How does it work?

- ▶ A **lock** is an object with two states (*open/closed*) and two operations (*to open/to close*)
- ▶ When creating the lock, it will be initially open.
- ▶ Usage of lock:

close lock

CS

open lock

Non-critical section



▶ Close lock:

- ▶ If open → it closes the lock
- ▶ If closed by another thread → the current thread waits
- ▶ If closed by itself → no effect

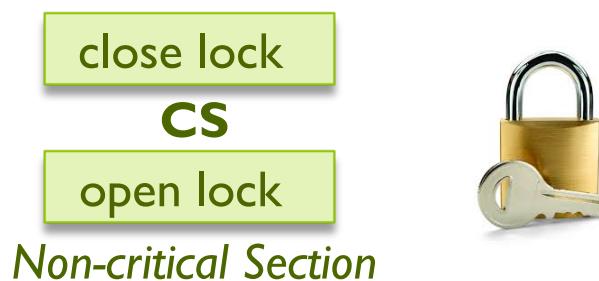
▶ Open lock:

- ▶ If already open → no effect
- ▶ If closed by another thread → no effect
- ▶ If closed by itself → the lock opens
 - ▶ If anyone is waiting, it gets closed by one of the threads that are waiting (who will keep on running)
 - The choice depends on the implementation, but it normally will meet equity



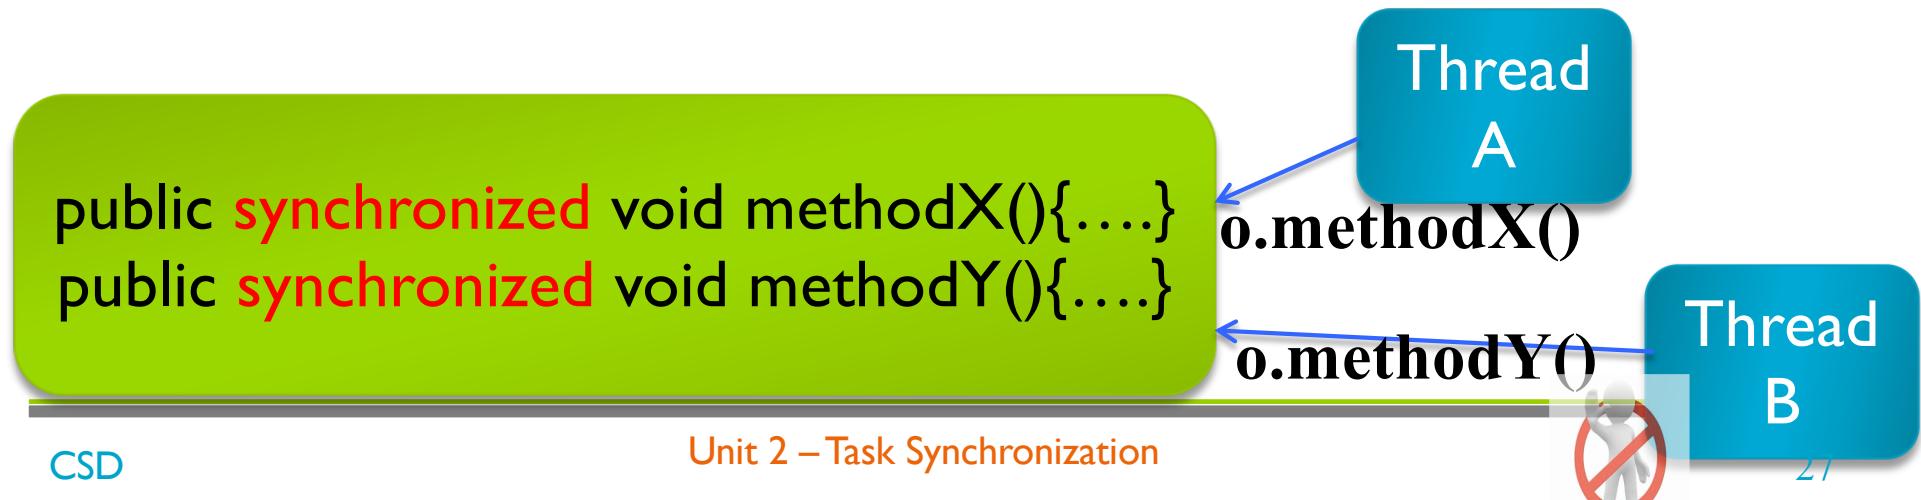
Locks → How does it work?

- ▶ Using locks we convert a critical section into an **atomic action**
 - ▶ **Atomicity:** only one thread can execute the protected code at any given time
 - ▶ Ensures reliable update of shared variables or shared objects
 - Free from race conditions or corruption of states
 - Only the steady states are visible
 - Ensures that every thread always access to the most recent value of each shared variable or shared object.



Locks in Java → How can you use them?

- ▶ Every object has an **implicit** associated lock
 - ▶ You must label with synchronized all methods that are part of the critical section
 - ▶ When Java enters a method labeled as ‘synchronized’, the lock gets closed, and it opens when leaving the method.
 - ▶ When a thread releases an intrinsic lock, a *happens-before* relationship is established with that action and any subsequent acquisition of the same lock.
 - ▶ For the same object, all its methods labeled with synchronized are executed in mutual exclusion among them





Locks in Java → How can you implement them?

- ▶ **Implementation** → add the **synchronized** label to every method that:
 - ▶ Modifies an attribute which another thread has to read later
 - ▶ Reads an attribute updated by another thread



Locks in Java → How can you implement them?

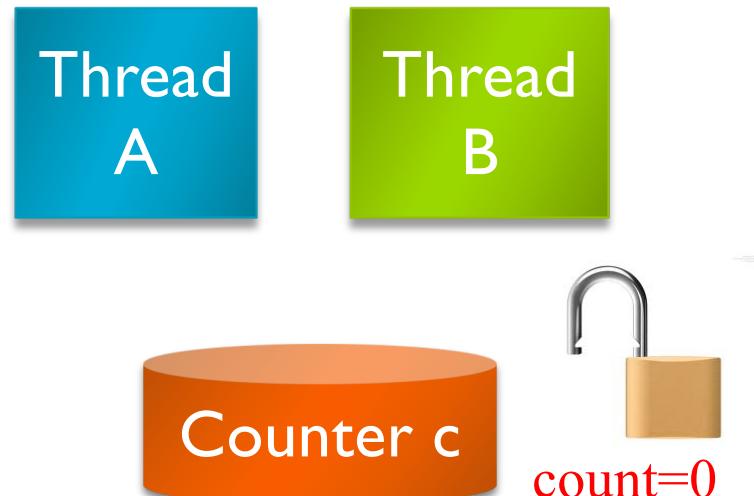
```
public class Counter {  
    private long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
  
    public synchronized long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter(); // initialized to 0
```

Example:
Counter Version
without race
conditions

Locks in Java → Let's see an example

► Interleaving example:

Initially: implicit lock of **Counter**
is open



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
  
    public synchronized long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter();
```

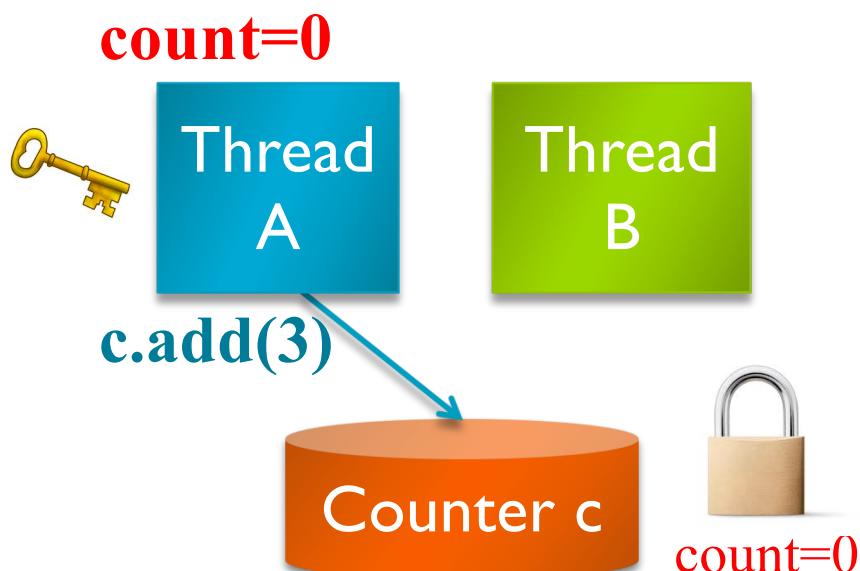
Locks in Java → Let's see an example

► Interleaving example:

Thread A executes `c.add(3)`

A loads `c.count` (local copy=0)

And Java closes the implicit lock



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter();
```

Locks in Java → Let's see an example

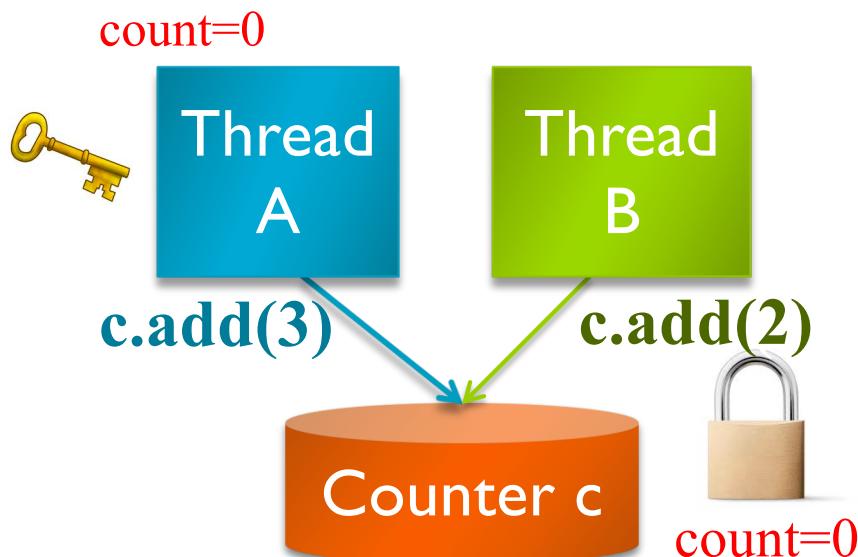
- ▶ Interleaving example:

Thread A executes c.add(3)

A loads c.count (local copy=0)

Thread B executes c.add(2)

B has to wait (lock already closed by another thread)



```
public class Counter {
    protected long count = 0;
    public synchronized void add(long x) {
        count += x;
    }
    public synchronized long getCount() {
        return count;
    }
    ...
}
```

Counter c= new Counter();

Locks in Java → Let's see an example

- ▶ Interleaving example:

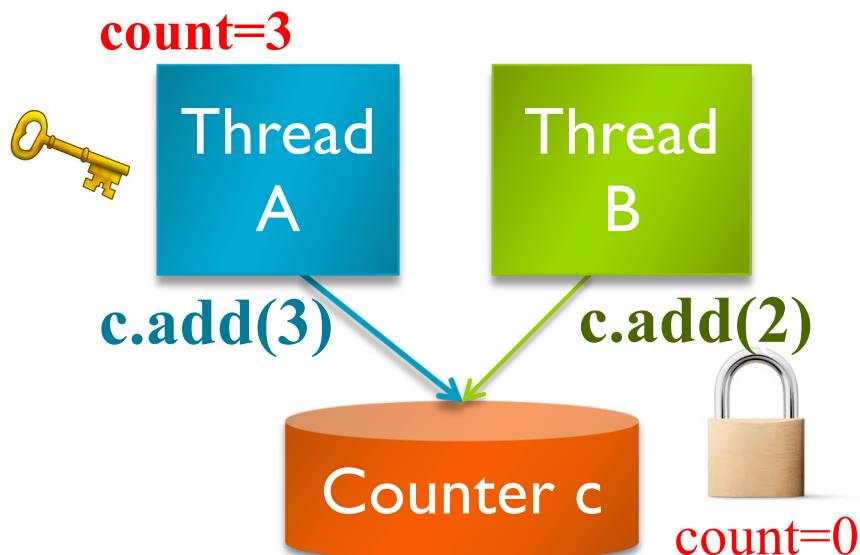
Thread A executes c.add(3)

A loads c.count (local copy=0)

A adds 3 (local copy=3)

Thread B executes c.add(2)

B has to wait (lock already closed by another thread)



```
public class Counter {
    protected long count = 0;
    public synchronized void add(long x) {
        count += x;
    }
    public synchronized long getCount() {
        return count;
    }
    ...
}
```

Counter c= new Counter();

Locks in Java → Let's see an example

► Interleaving example:

Thread A executes `c.add(3)`

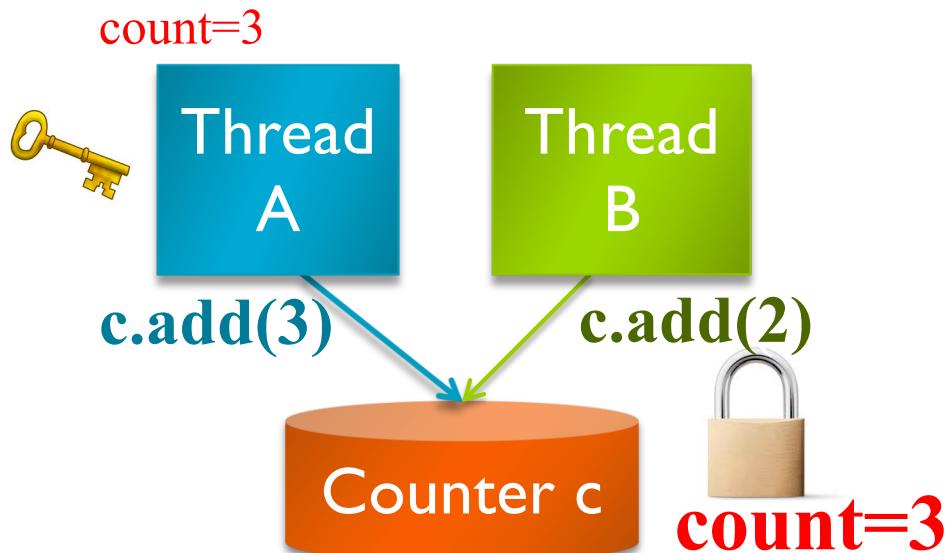
A loads `c.count` (local copy=0)

A adds 3 (local copy=3)

A updates the heap (`c.count=3`)

Thread B executes `c.add(2)`

B has to wait (lock already closed by another thread)



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Locks in Java → Let's see an example

► Interleaving example:

Thread A executes c.add(3)

A loads c.count (local copy=0)

A adds 3 (local copy=3)

A updates the heap (c.count=3)

A leaves the synchronized method (and Java opens the lock)

count=3

Thread
A

c.add(3)

Thread
B

c.add(2)

Counter c

count=3

Thread B executes c.add(2)

B has to wait (lock already closed by another thread)

```
public class Counter {
    protected long count = 0;
    public synchronized void add(long x) {
        count += x;
    }
    public synchronized long getCount() {
        return count;
    }
    ...
}
```

Counter c= new Counter();

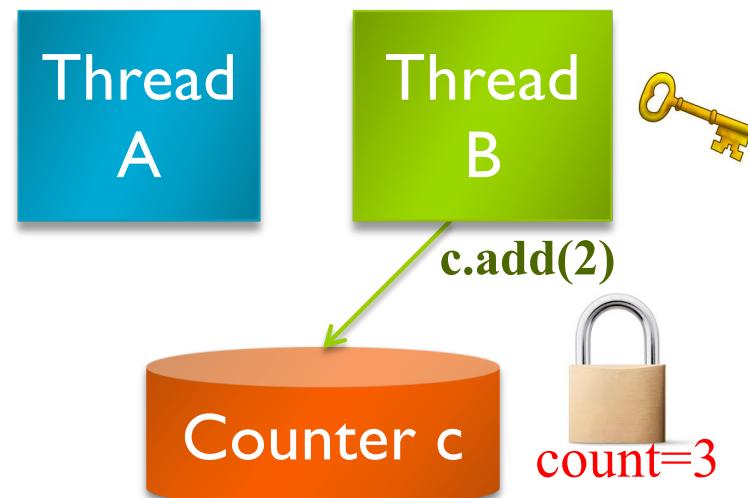
Locks in Java → Let's see an example

► Interleaving example:

Thread A executes `c.add(3)`

Thread B executes `c.add(2)`

B awakes and Java closes the lock



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
    ...  
    Counter c = new Counter();
```

Locks in Java → Let's see an example

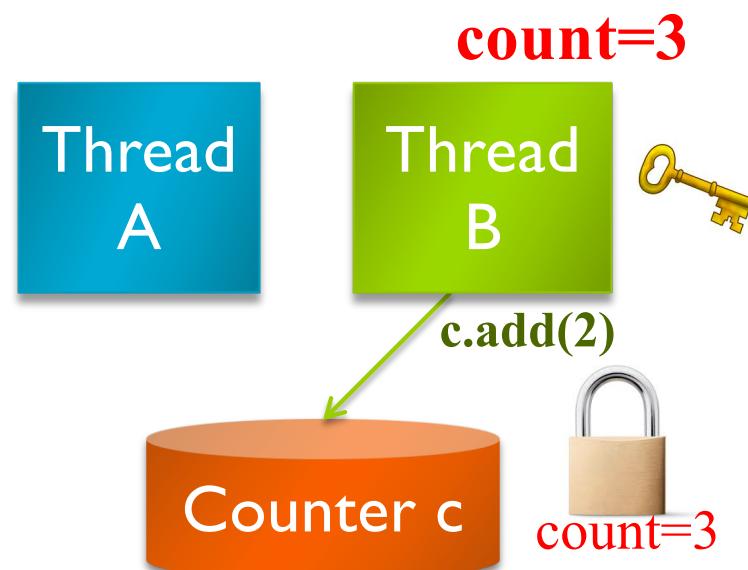
► Interleaving example:

Thread A executes `c.add(3)`

Thread B executes `c.add(2)`

B awakes and Java closes the lock

B loads c.count (local copy=3)



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
  
    public synchronized long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter();
```

Locks in Java → Let's see an example

► Interleaving example:

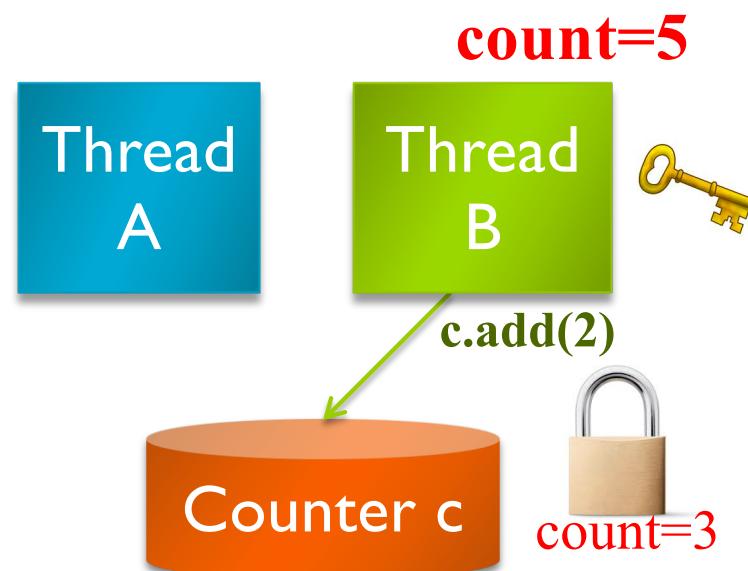
Thread A executes `c.add(3)`

Thread B executes `c.add(2)`

B awakes and Java closes the lock

B loads `c.count` (local copy=3)

B adds 2 (local copy=5)



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Locks in Java → Let's see an example

► Interleaving example:

Thread A executes `c.add(3)`

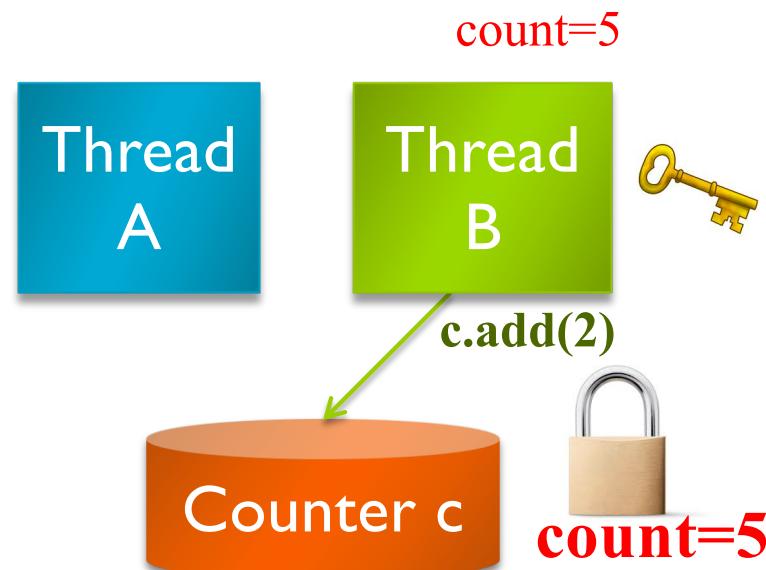
Thread B executes `c.add(2)`

B awakes and Java closes the lock

B loads `c.count` (local copy=3)

B adds 2 (local copy=5)

B updates the heap (`c.count=5`)

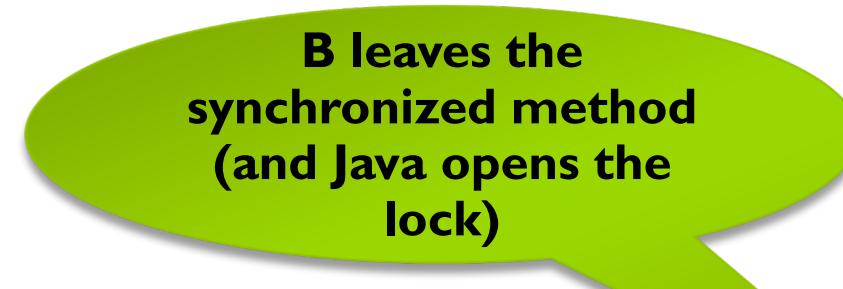


```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Locks in Java → Let's see an example

► Interleaving example:

Thread A executes `c.add(3)`



Thread B executes `c.add(2)`

B awakes and Java closes the lock
B loads `c.count` (local copy=3)
B adds 2 (local copy=5)
B updates the heap (`c.count=5`)

```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
    public synchronized long getCount() {  
        return count;  
    }  
    ...  
    Counter c= new Counter();
```

Locks in Java → Let's see an example

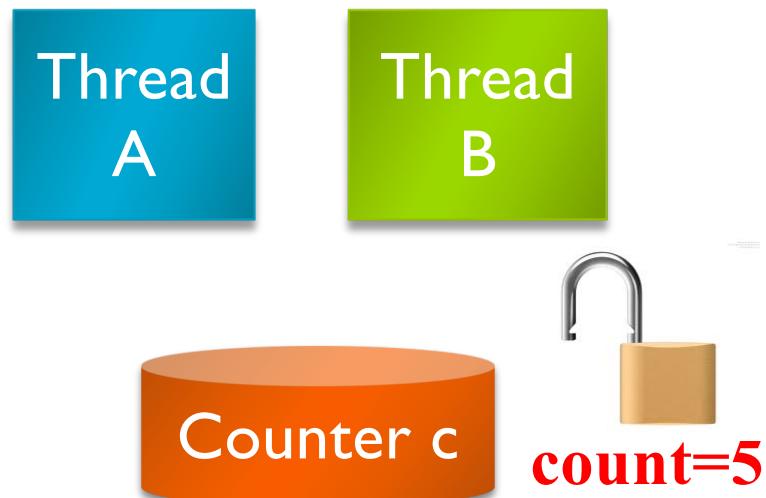
- ▶ Independent of interleaving:

Thread A executes `c.add(3)`

Final value: ALWAYS 5!!

Thread B executes `c.add(2)`

Code free of **race conditions**



```
public class Counter {  
    protected long count = 0;  
    public synchronized void add(long x) {  
        count += x;  
    }  
  
    public synchronized long getCount() {  
        return count;  
    }  
}  
...  
Counter c= new Counter();
```



Java synchronization

- ▶ Java provides two basic synchronization idioms:
 - ▶ **Synchronized methods:** using **synchronized** label at method declaration.
 - ▶ **Synchronized statements:** specifying the object that provides the intrinsic lock
- ▶ Allows using more than one lock inside the same method.

- Methods *inc1* & *inc2* can interleave
- But updating **c1** is done in mutual exclusion
- Same for **c2**

```
public class MsLunch{  
    private long c1 = 0;  
    private long c2 = 0;  
    private Object lock1 = new Object();  
    private Object lock2 = new Object();  
  
    public void inc1() {  
        synchronized(lock1){ c1++; }  
    }  
  
    public void inc2() {  
        synchronized(lock2){ c2++; }  
    }  
}
```



Conditional Synchronization

- ▶ The concept of **critical section** prevents interferences when accessing to shared variables
- ▶ But we must also solve the problem of *conditional synchronization*:
 - ▶ Make a thread wait until certain condition is met
 - ▶ In next unit we will study the concept of *monitor*
 - ▶ This solves both problems (i.e. mutual exclusion and conditional synchronization) using a single construction



Learning results of this Teaching Unit

- ▶ At the end of this unit, the student should be able to:
 - ▶ Identify the sections of an application that can be executed concurrently by different activities.
 - ▶ Identify the problems related with communication using shared memory. Describe the determinism problem.
 - ▶ Identify the different types of synchronization.
 - ▶ Distinguish the critical sections of an application and protect them using the suitable mechanisms.